SPARLI AMAN

Senior AI Engineer · Agentic Systems, Integrations & Data Pipelines

Islamabad, Pakistan (UTC+5) | Remote or relocation
hello@sparli.dev | sparli.dev | github.com/Sparli | linkedin.com/in/sparli-aman

Seven years shipping production software, the last four building the AI and data layer of **Shurq**, a UK Amazon-analytics SaaS I joined at the idea stage and helped grow into a paid product with a 10+ person team and 200+ concurrent users. I build agents that take real actions in business systems (ad platforms, CRMs, telephony) behind approval gates and audit logs, and the data pipelines they run on. I've delivered end to end for clients in the US, Canada, the UK and the Gulf, and I want to go deep on one product with one team.

EXPERIENCE

Shurq · Founding AI & Data Engineer, leading AI and data engineering

Jul 2022 – present

UK, remote · Amazon seller analytics and PPC automation SaaS (shurq.com) · joined at idea stage

Joined as the first AI/data engineer. Shurq is now a paid SaaS (v2.76, 670+ shipped releases) serving Amazon brands across **21 marketplaces**, run by a 10+ person team, with **200+ concurrent users** on the platform.

Data pipeline (owner). Built and run the pipeline that pulls sales, fee and advertising data from Amazon every hour and turns it into true per-product profit across all 21 marketplaces. Every dashboard, alert and AI agent in the product reads from it.

MCP server (led and built). Shurq's MCP server exposes **1,080 tools, 29 of them write actions**, over streamable HTTP with scoped API keys, so a seller can query and act on their live Amazon account from Claude, ChatGPT, Cursor or Claude Code. It is also the tool layer behind the product's AI "employees", which act on their own only after earning a trust score.

AI Ad Manager (led and built). Rules, scheduled workflows and an approval queue for bid, budget and keyword changes. Every change is logged and reversible, and nothing writes to a live ad account without a rule or a human releasing it. The LLM layer falls back from self-hosted Ollama to OpenAI, so a provider outage costs latency, not the feature.

Python · TypeScript · PostgreSQL · BigQuery · MCP · Amazon Ads & Selling Partner APIs · OpenAI · Ollama · AWS

Croncore · Founder & Lead Engineer, AI engineering consultancy

Jan 2025 – present

Scope, build, deploy and support AI systems for clients in the US, Canada, Saudi Arabia and Mongolia, usually as the sole engineer. Selected work below.

Freelance AI / Full-Stack Engineer · Top Rated on Upwork

2019 – 2022

SELECTED CLIENT & PRODUCT WORK

EduImmersion: AI learning platform + school-network dashboards

2025 – present · in pilot

Québec EdTech client · took raw requirements to two shipped products

A **learner platform** that generates lessons, quizzes and reflections by grade, language and pace (free tier plus US$29.90/month Premium on Stripe). And a **role-based dashboard suite** for school service centres, schools, teachers and parents. Both are EN/FR, with data hosted in Canada, now in pilot in Québec.

Built the lesson-image pipeline. A rework of its prompts and text handling **cut text-rendering errors by ~80%** across 8 EN/FR test scenarios, with no model change, cost or latency added. An EN/FR-parity Playwright suite blocks a release when either language regresses. Wrote the spec for the next stage: content, assessment and analytics agents sharing one learner record in Postgres.

Next.js 16 / React 19 · Django/DRF · PostgreSQL · Celery · Redis · Gemini · Stripe · Playwright

Voice agents that act in the CRM mid-call

2026

US payment-services client (sole engineer) · Canadian SEO SaaS

Outbound sales-calling platform: webhook server, post-call processor, Pipedrive dedupe, deal creation and callback scheduling, 10-scenario follow-up email routing, number-pool rotation, and an SMS tool that sends a document **during the call**. Compliance built in: timezone dialling windows, DNC scrubbing, mid-call opt-out detection, verbal-consent logging, CASL unsubscribes.

Tuned on live calls: switched the conversational model to `claude-haiku-4-5` to fix response latency, and traced early hang-ups to speech-to-text hearing "hello" as "no". The fix was a never-end-on-first-turn rule.

Node.js · Twilio · ElevenLabs Conversational AI · Synthflow · Pipedrive · Google Places · Render

Arabic (Saudi-dialect) inbound sales agent, live on web and phone 2026

Answers in about 0.5 s by browser call or phone number. It captures the lead and hands off to a human. Model fallback (gpt-oss-120b to Llama 3.3 on Groq) plus a scripted fallback means a live demo never goes silent, and Twilio webhooks are signature-verified.

Next.js · Vercel AI SDK · Groq · Whisper · Twilio

Agents with write access to business systems, made safe 2026

leadflow-agent (public, github.com/Sparli): turns an inbound enquiry into HubSpot records. A human approves each write, every action goes to a ledger with one-command undo, and it costs under $0.002 per lead. CI runs strict mypy and gitleaks. **Croncore delivery agent:** one tool surface over the Anthropic API and the Claude Agent SDK. It triages GitHub PRs, tracks ClickUp tasks and drafts client reports. Outbound actions wait in a `pending_actions` table for release, and every run is scoped to one project so data can't cross between clients.

Python · Anthropic API · Claude Agent SDK · Next.js · Drizzle/Postgres · HubSpot, ClickUp, GitHub, Gmail APIs

Also built 2025 – 2026

Speech service into an existing phone system (Mongolian client): FastAPI service with two speech engines behind one interface, async jobs with retry and webhook callbacks, and audio matched to the client's Asterisk/FreePBX format so nothing on their side changed. **Offline-first retail POS** (Tauri/Rust + React): ~380 automated tests, encrypted backups, money held as integers, never floats. **Live speech-to-speech translation**: Deepgram, DeepL, ElevenLabs, with latency shown on screen. **MCP coordination server** that lets coding agents see each other's file claims. **Prospect-research pipeline** on a VPS: free models, strict quality gates, nothing sent without review. **Multi-user time tracker** on Cloudflare Workers/D1. Led a retrieval agent that quotes live prices from an automotive client's own sheets (pgvector).

SKILLS

Languages	Python · TypeScript/JavaScript · Node.js · SQL · Rust (Tauri)
Agentic & LLM	MCP servers · tool-using agents · Anthropic API & Claude Agent SDK · OpenAI · Gemini · Groq · Vercel AI SDK · structured outputs · RAG (pgvector) · provider fallback · human-in-the-loop approval gates · prompt-injection containment · eval design
Voice	ElevenLabs Conversational AI · Twilio voice/SMS · Deepgram · Whisper · Synthflow · Vapi · Asterisk/FreePBX
Backend & data	FastAPI · Django/DRF · Next.js · PostgreSQL · BigQuery · Redis · Celery · Drizzle · SQLite · scheduled data pipelines · rate limiting and backoff
Cloud & ops	AWS · GCP · Vercel · Render · Cloudflare Workers/D1 · Docker · GitHub Actions CI · Playwright · pytest
Integrated with	Amazon Ads & SP-API · HubSpot · Pipedrive · ClickUp · GitHub · Gmail · Google Sheets/Places · Stripe

EDUCATION

BSc Software Engineering, Pak-Austria Fachhochschule, Institute of Applied Sciences and Technology, Pakistan